Personalized Foundation Model for Social Simulation

Wei, Zhongyu

Fudan University | Shanghai Innovation Institute

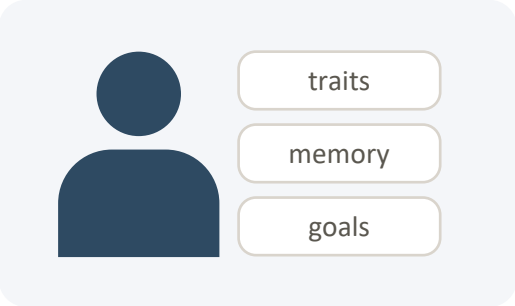
2026.10.09

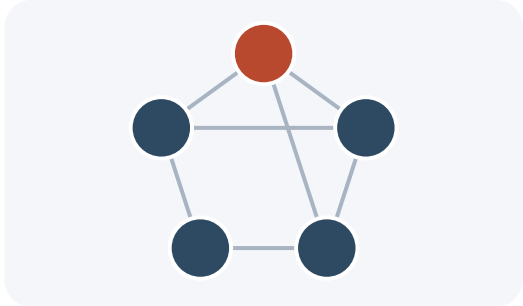
Second Workshop on Social Simulation with LLMS: Fidelity in Applications

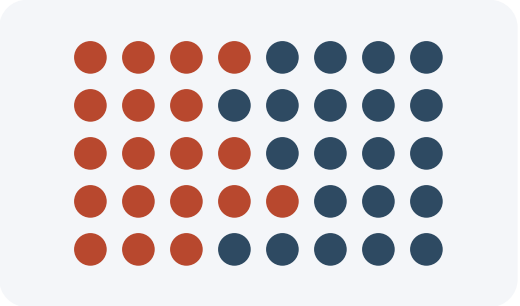
@ COLM 2026, Hilton Union Square, San Francisco, CA, USA

What Is Social Simulation?

Social simulation starts with individuals. We model how they behave, allow them to interact, and observe the collective patterns that emerge.







01 INDIVIDUALS

Model how they behave

02 INTERACTIONS

Let them interact

03 COLLECTIVE PATTERNS

Observe what emerges

DOWNSTREAM APPLICATIONS

Social Science

Test how individual behavior gives rise to collective phenomena

Policy Evaluation

Try interventions in a simulated society before real-world roll-out

Market Analysis

Anticipate how consumers and markets respond

Behavior as a function of person and environment

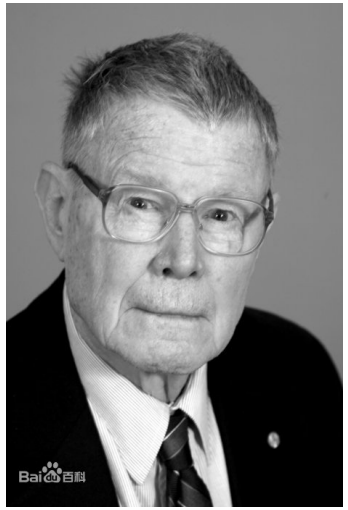
Kurt Lewin's Field Theory (1936)

$$B = f(P, E)$$

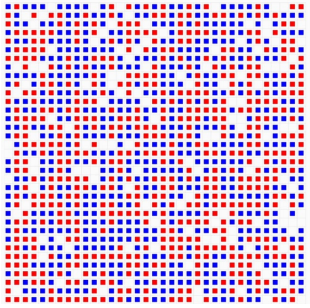
B	f	P	E
BEHAVIOR	**FUNCTION**	**PERSON**	**ENVIRONMENT**
What people do	**How they decide**	**Who they are**	**What they face**

Lewin, field theory, 1936

Agent-based Modeling: Schelling's segregation model (1971)

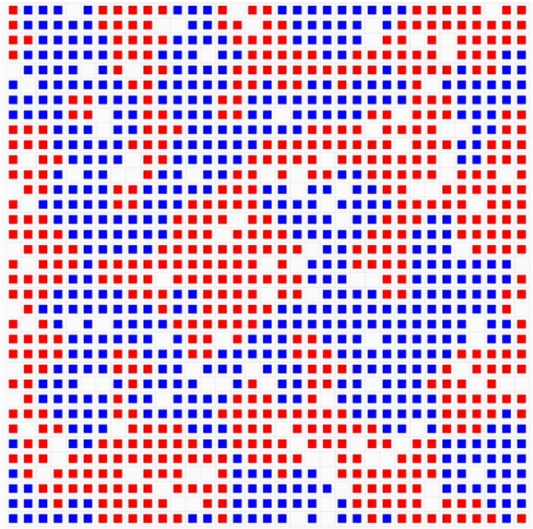


Thomas Crombie Schelling
2005 Nobel Laureate



Initial state

- **Goal:** How individual preferences shape segregation
- **Action space:** move / stay (by neighbor similarity)

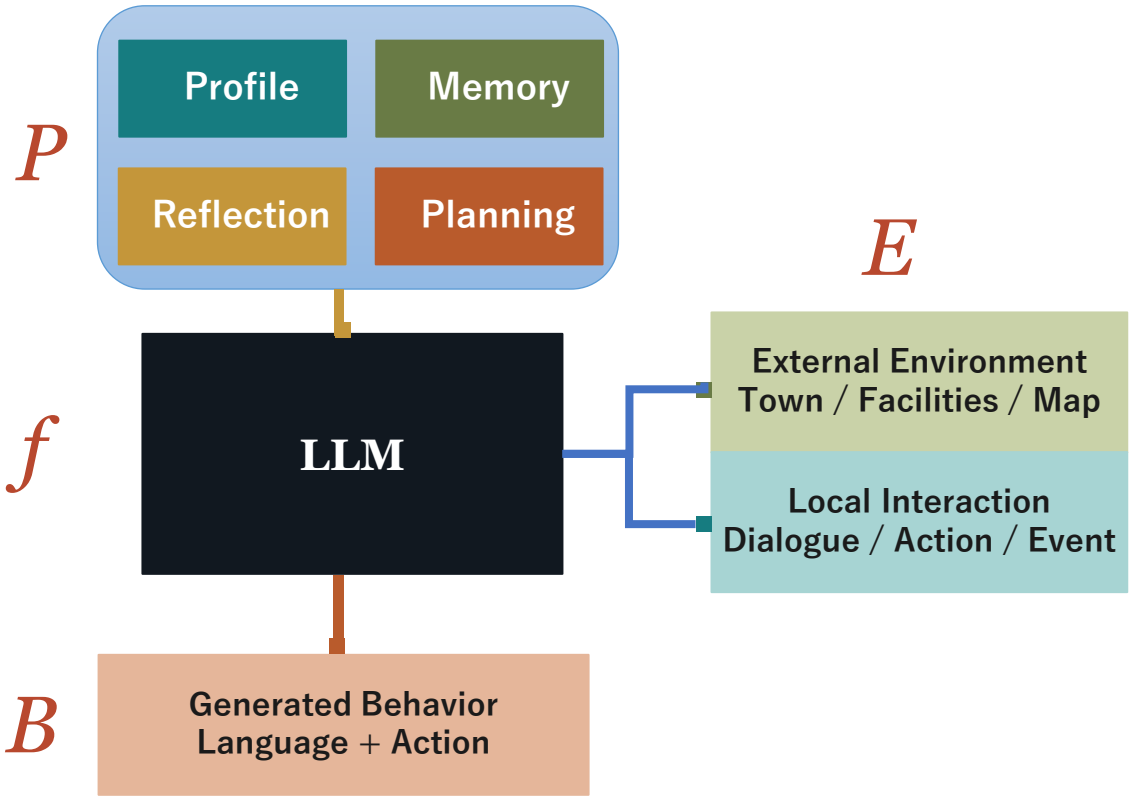


Threshold
30%

P	PERSON Single-attribute agent: one category label.
E	ENVIRONMEN Local information from the eight surrounding neighbors.
f	FUNCTION Threshold rule: move if similar neighbors < X%.

ABM Unleashed: Generative Agents





Stanford Generative Agents

[Joon S. Park et al., 2023]

The first **LLM-driven** multi-agent sandbox

Expanding the **scope of social simulation**

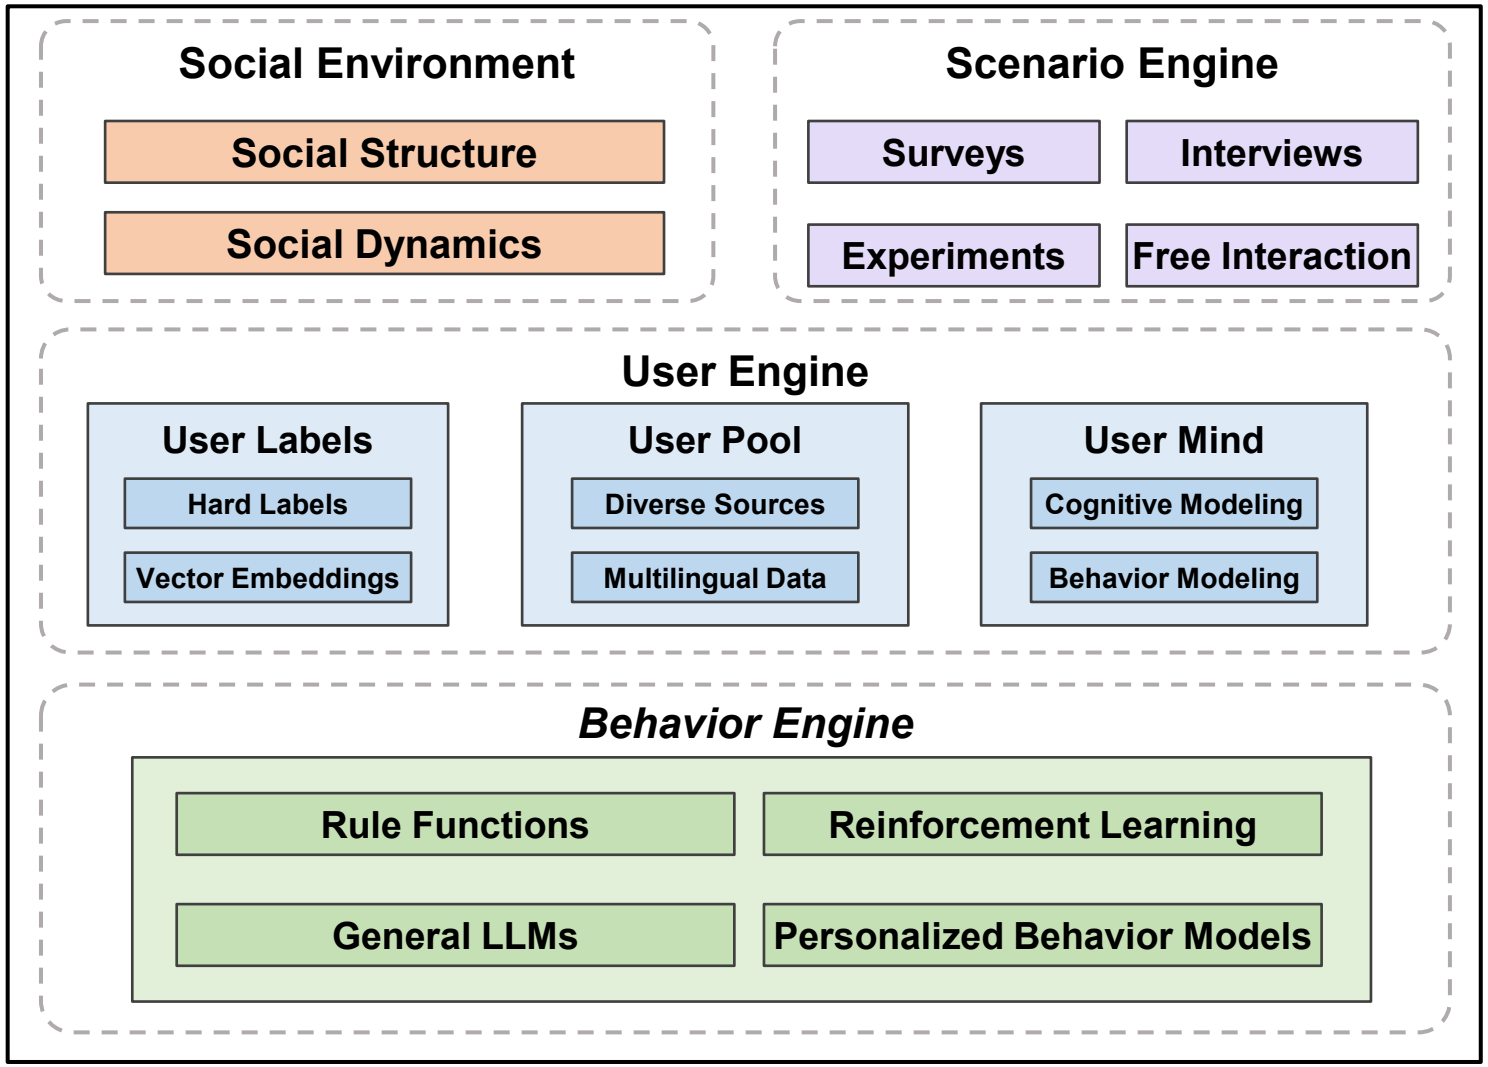
SocioVerse: A World Model for Social Simulation

$$B = f(P, E)$$

10M syn-user pool

Automated labeling

3 application cases



Socioverse: A world model for social simulation powered by LLM agents and a pool of 10 million real-world users, Arxiv 2025

HiSim: Low-Cost Simulation with Hybrid Agents

Behavior Engine: Hybrid Behavior Engine

LLM -> core users, ABM -> ordinary users

Scenario Engine: Digital-twin Social Platform

Public & private information channels, real-time clock

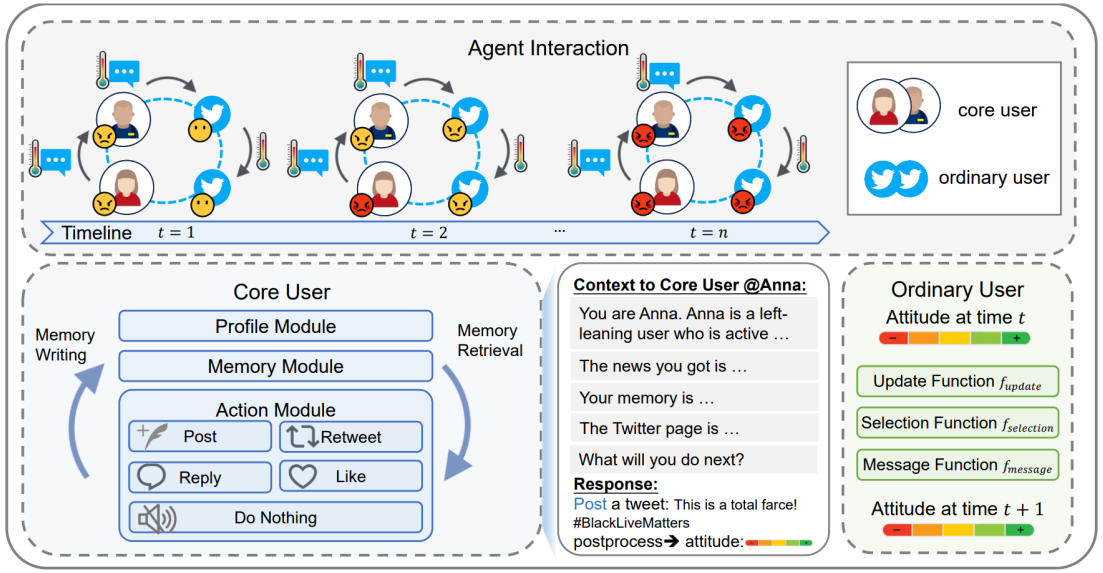
User Engine: Driven by Real-user Replicas

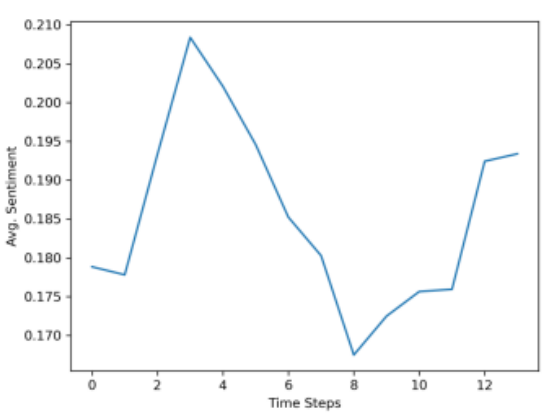
1,000 target users, 14-day behavior prediction

Alignment: Real-event Benchmarks

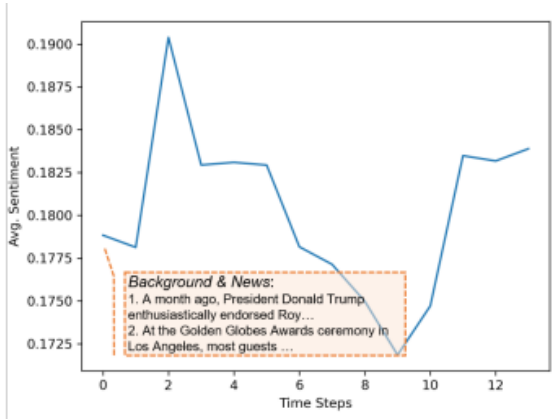
#metoo, #blacklivematters, #roeovertuned

Sentiment distribution of public opinion





Real Attitudes

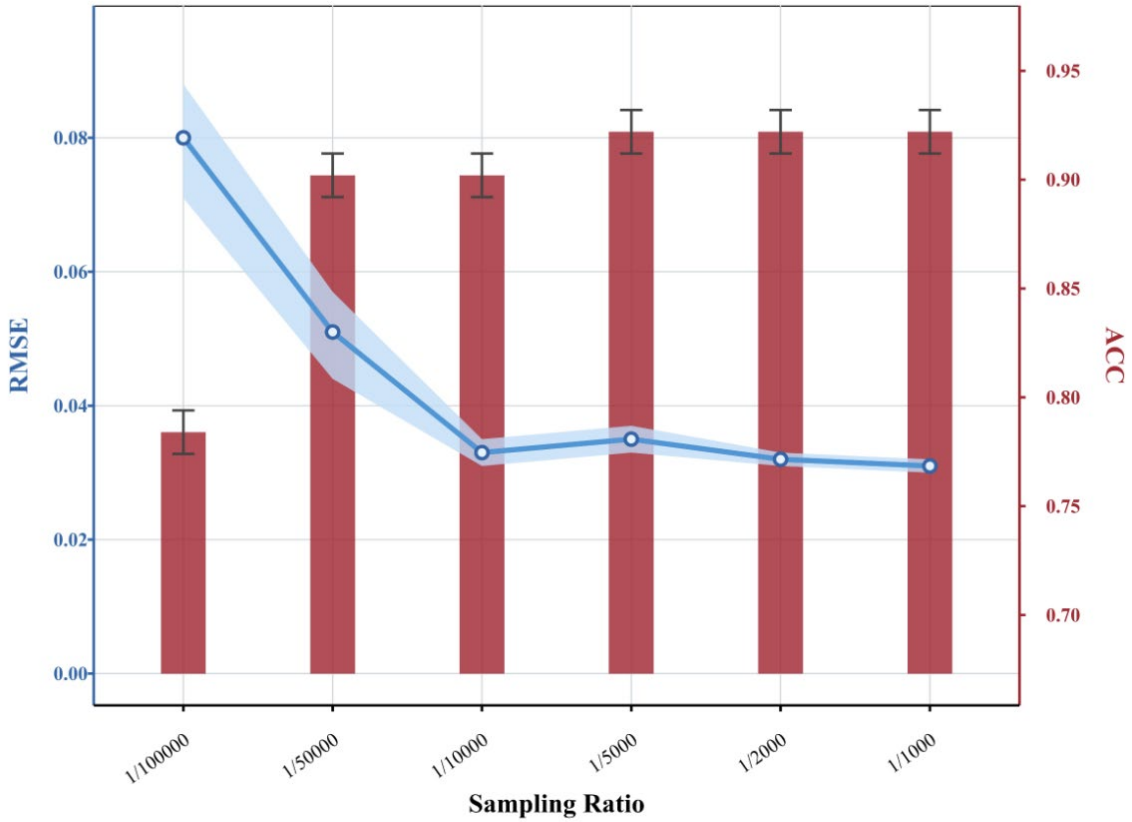


Simulated Attitudes

Unveiling the truth and facilitating change: Towards agent-based large-scale social movement simulation, ACL 2024

ElectionSim: Massive Population Election Simulation

- **User Engine**
  - Million-size user pool
- **Scenario Engine**
  - 1 VS 1 questionnaire from ANES
- **Behavior Engine**
  - General LLM + prompt
- **Social Environment**
  - Human crafted social events
  - Candidate information



- **Evaluation results: ~90% on individual level, <4 % on state level**

ElectionSim: Massive Population Election Simulation Powered by Large Language Model Driven Agents, Arxiv 2024

ConsumerSim: Event-driven Confidence Forecasting

User Engine: Real-to-Virtual Hybrid: Population built from **real survey data**

Environment Engine: Socio-economic statistics + matched news

Alignment: Consumer Confidence Indices (US, EU, Japan)

Forecast accuracy reaches 95%, ahead of prior methods (75 months back evaluation)

Target / method	MAE ↓			RMSE ↓			Pearson r ↑			Spearman ρ ↑		
	U.S.	EU27	Japan	U.S.	EU27	Japan	U.S.	EU27	Japan	U.S.	EU27	Japan
ConsumerSim	**3.45**	**1.607**	**1.363**	**4.56**	**2.737**	**2.040**	**0.9590**	**0.882**	**0.824**	**0.9519**	**0.849**	**0.840**
AR Ridge	3.77	2.152	1.393	5.18	4.205	2.295	0.9459	0.749	0.794	0.9363	0.806	0.839
Expectations Reg.	3.96	2.490	1.643	5.45	4.082	2.774	0.9443	0.834	0.760	0.9346	0.799	0.804
News Semantic FAAR	4.02	2.380	1.550	5.63	3.769	2.374	0.9378	0.791	0.753	0.9331	0.797	0.816
3-Month Mean	4.34	2.216	1.816	5.88	3.422	2.722	0.9295	0.804	0.652	0.9092	0.782	0.729
SARIMAX Macro	4.93	3.072	2.854	7.98	5.191	4.932	0.8933	0.755	0.609	0.9082	0.735	0.648
ETS Damped Trend	4.96	1.833	1.492	6.94	3.471	2.438	0.9058	0.830	0.798	0.8767	0.833	0.833
Theta Model	5.08	1.681	1.374	6.95	2.848	2.063	0.9030	0.870	0.818	0.8775	0.840	0.830
Auto ARIMA	5.44	2.373	1.642	7.76	4.173	2.576	0.8852	0.767	0.778	0.8707	0.775	0.775
Prophet Macro	5.89	4.148	2.629	8.10	5.664	3.664	0.8985	0.564	0.541	0.8965	0.384	0.596
Prophet Univ.	6.06	4.806	2.801	7.96	6.175	3.596	0.8714	0.274	0.468	0.8277	0.321	0.555
Macro-Market-News Ridge	6.21	2.511	1.951	14.40	3.810	2.841	0.7215	0.836	0.783	0.8758	0.780	0.763
Macro Stable Ridge	6.45	2.490	1.959	15.12	4.082	2.997	0.6893	0.834	0.771	0.8705	0.799	0.783

Uncovering Salience-Driven Dynamics in Consumer Confidence with Generative Social Simulation, ArXiv 2026

How the Behavior Function Works Today

P PERSON CONTEXT

Demographics · personality · values
Past behavior and memories

E ENVIRONMENT CONTEXT

Scenario and task (survey item, platform)
Events, news and social feed

f

General-purpose LLM

prompted with P + E,
not trained for individuals

B BEHAVIOR

Survey answers · posts
Choices and actions

A prompt in practice: *"You are a 45-year-old teacher in Ohio, moderately conservative …"* + *"This week's news: … Question: How much do you trust the federal government?"* → ***"Not very much."***

Is a general-purpose LLM enough as the behavior engine?

SocioBench – A Benchmark for Evaluating LLMs' Capacity in Social Surveys

10 survey domains

5,000 respondents

30+ countries

600+ questions


1 Citizenship


2 Environment


3 Family & Gender


4 Health & care


5 National Identity


6 Religion


7 Government Role


8 Social Inequality

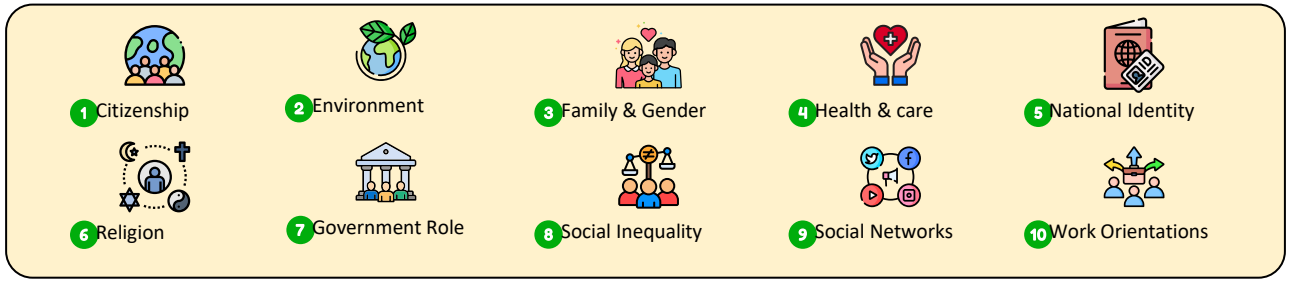

9 Social Networks


10 Work Orientations

SocioBench: Modeling Human Behavior in Sociological Surveys with Large Language Models, EMNLP 2025

SocioBench – A Benchmark for Evaluating LLMs' Capacity in Social Surveys

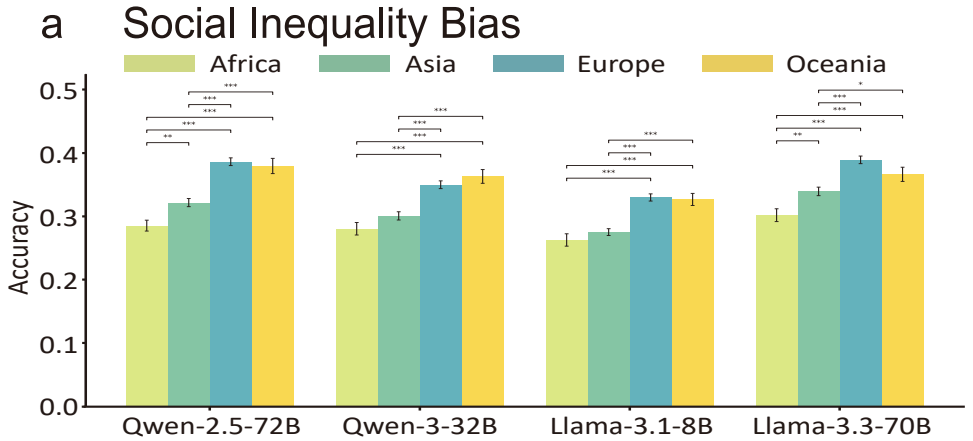
Built on ISSP – covering 10 domains, 5,000 respondents, and over 600 questions.

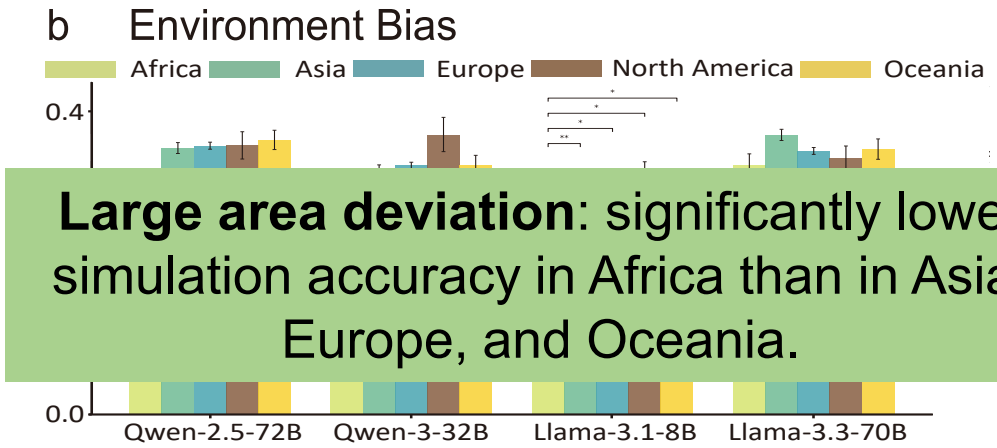


Model	Citizen	Enviro	Family	Health	Nat.Ident	Religion	R.Gov	S.Ineq	S.Net	Work	Average
					Accuracy % (↑)						
I										35.56	34.48
G										37.48	35.36
G										37.24	35.14
Y										34.73	34.06
M										35.33	33.21
Mixtral-8x22B-Instruct-v0.1	40.68	33.31	34.50	33.31	33.41	38.01	35.20	31.95	34.07	35.42	34.99
Llama-3.1-8B-Instruct	38.32	30.49	32.54	32.90	33.69	38.00	34.21	30.00	34.45	33.16	33.77
Llama-3.3-70B-Instruct	41.35	35.16	**38.51**	**36.83**	**38.35**	**41.48**	**39.11**	35.33	35.67	37.99	**37.98**
Q										31.79	32.95
Q										35.76	35.41
Q										36.58	36.75
Q										34.25	34.42
Q										34.80	35.10
D										35.69	35.23
D										**39.32**	37.92

In complex scenarios, individual behavior simulation accuracy is below 40%.

Large domain deviation: significantly lower accuracy in domains such as social inequality and environmental issues.





Large area deviation: significantly lower simulation accuracy in Africa than in Asia, Europe, and Oceania.

The social simulation bottleneck

① DATA

Distribution collapse

Model Collapse · Nature 2024

Synthetic training data can lose **long-tail signals**.

② TRAINING

Alignment tax

GPT-4 Technical Report

Uniform RLHF may **flatten and harm** some capabilities.

③ INFERENCE

Simulation bias

OpinionQA · ICML 2023 (Oral)

General LLMs still struggle to **simulate populations**.

Social simulation calls for foundation models built specifically to simulate individuals.

Bisbee et al., Political Analysis 2024; Goel et al., 2025; Samuel et al., 2024.

Individual Simulation: Evolving Paradigms

1 Prompt Conditioning

e.g., Out of One, Many (Argyle et al., 2023)

2 Retrieval & Memory

e.g., Generative Agents (Park et al., 2023); LaMP (Salemi et al., 2024)

3 Parametric Personalization

e.g., OPPU (Tan et al., 2024)

4 Behavior Foundation Model

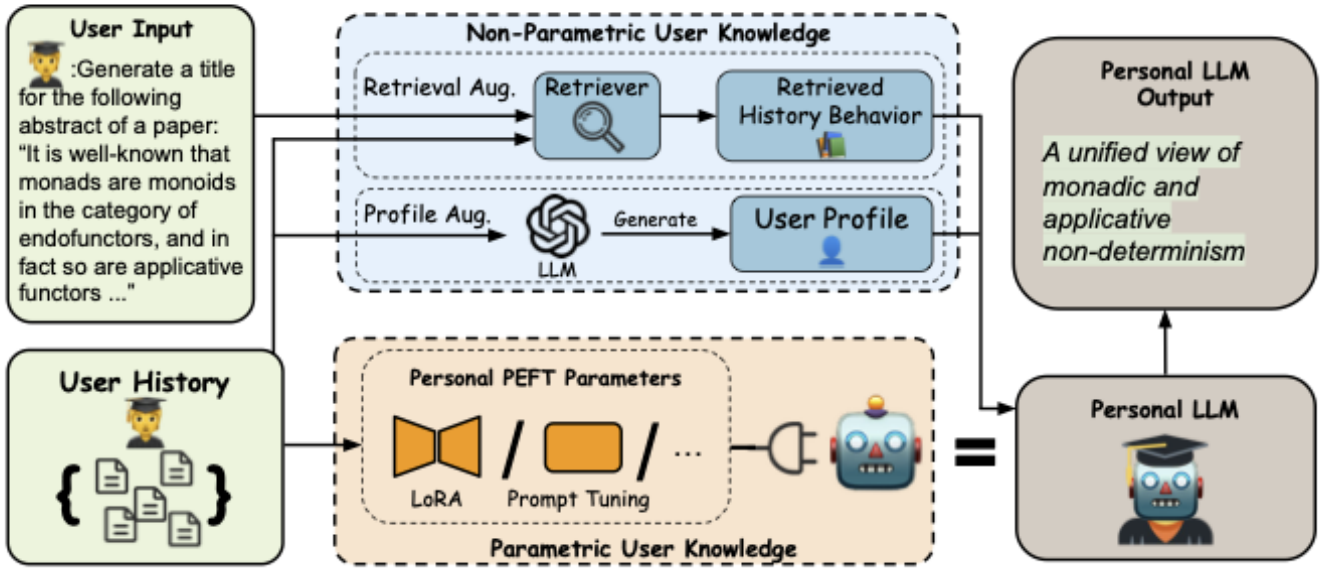
e.g., HumanLLM (Lei et al., 2026); OdysSim (Zhou et al., 2026)

2023– 2024– 2025– 2026–

Personalized LLMs are evolving from adapting to individuals toward understanding and simulating human behavior.

Argyle et al., 2023; Park et al., 2023; Salemi et al., 2024; Tan et al., 2024; Lei et al., 2026; Zhou et al., 2026.

Parametric Representation



Approach	Per-user storage	Main tradeoff
Encoder / soft prompt	Low	Compression
Personal PEFT	Grows with users	Training and storage
Shared modules / router	Controllable	Module entanglement

Tan et al., OPPU, EMNLP 2024.

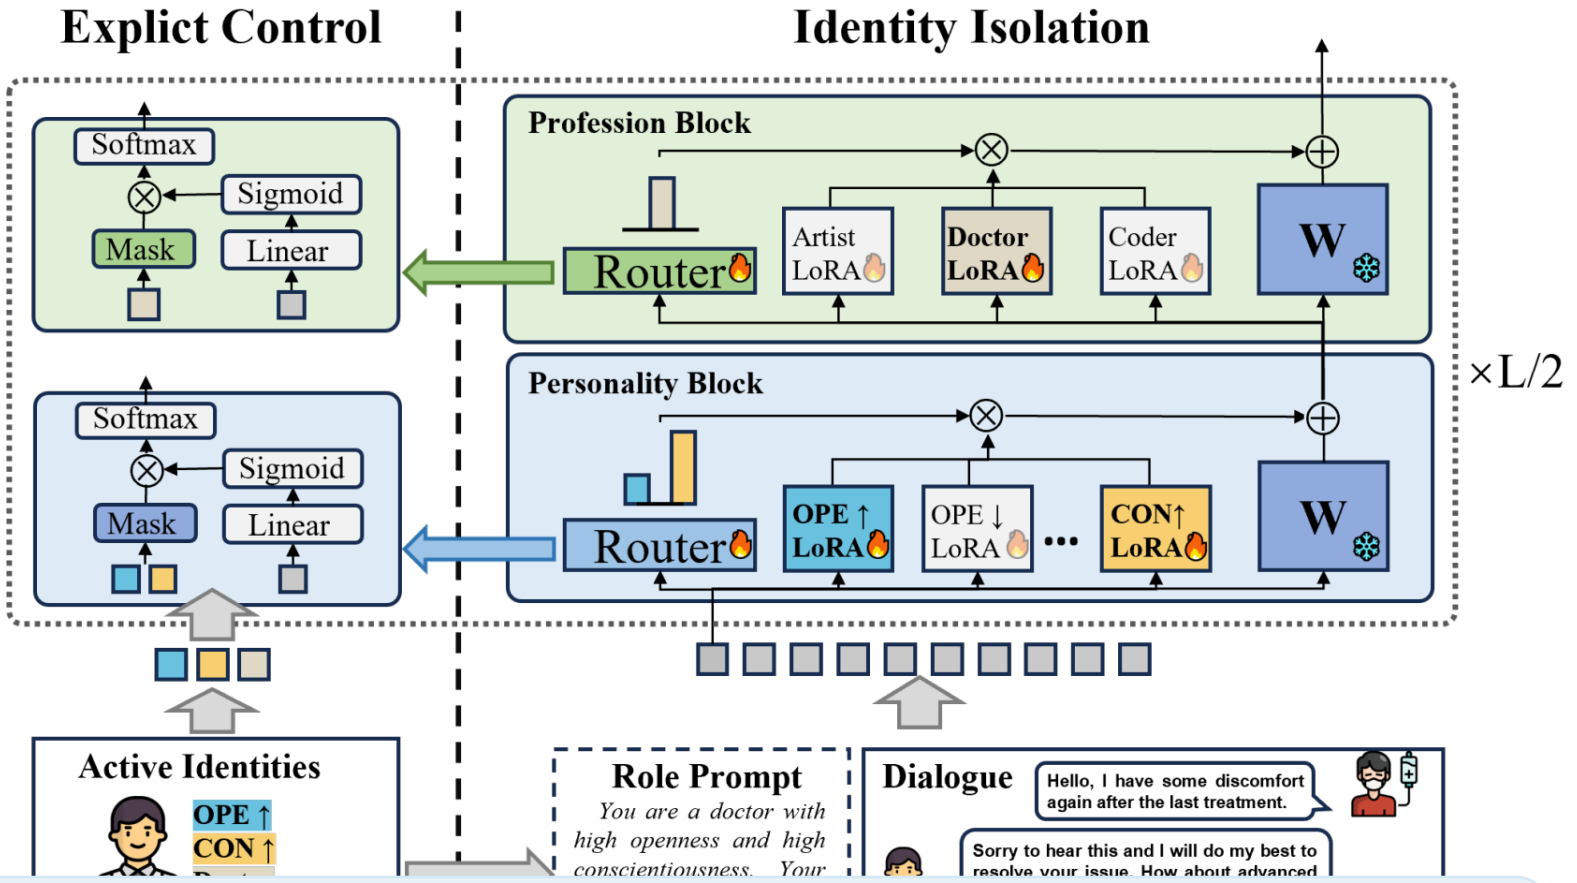
HIRPF: Identity-Specific Parameters

Identity isolation

Identity-specific LoRA blocks across model layers.

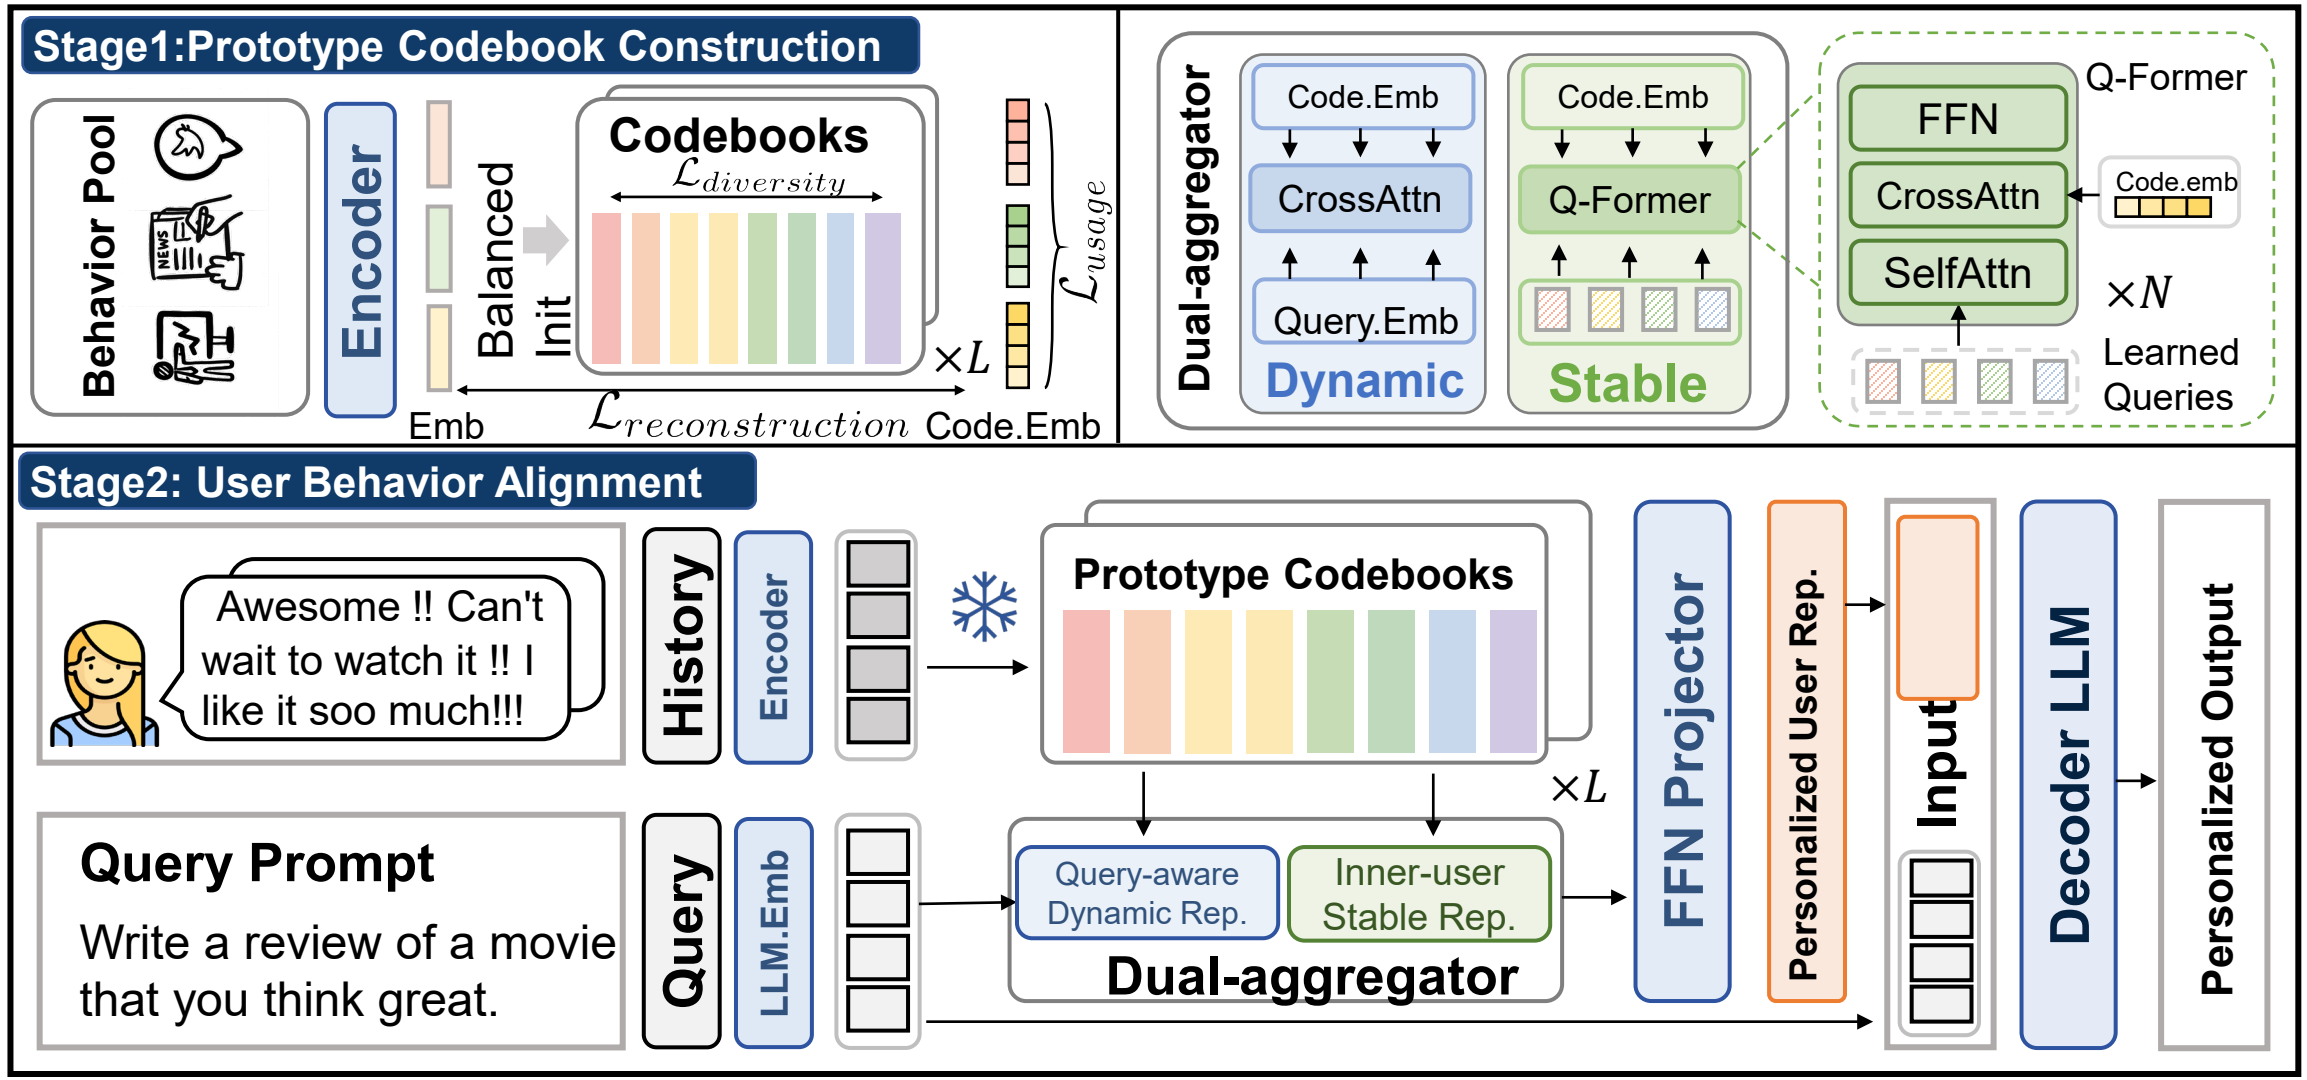
Parameter activation

Activate and update only identity-specific parameters.



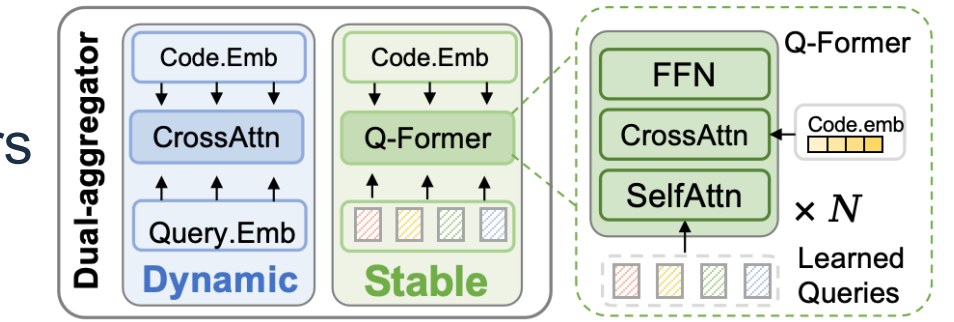
Limitation: Predefined set of identity labels captures only part of an individual's complexity.

CURP: Codebook-based User Representation



A person has a stable core and a dynamic state

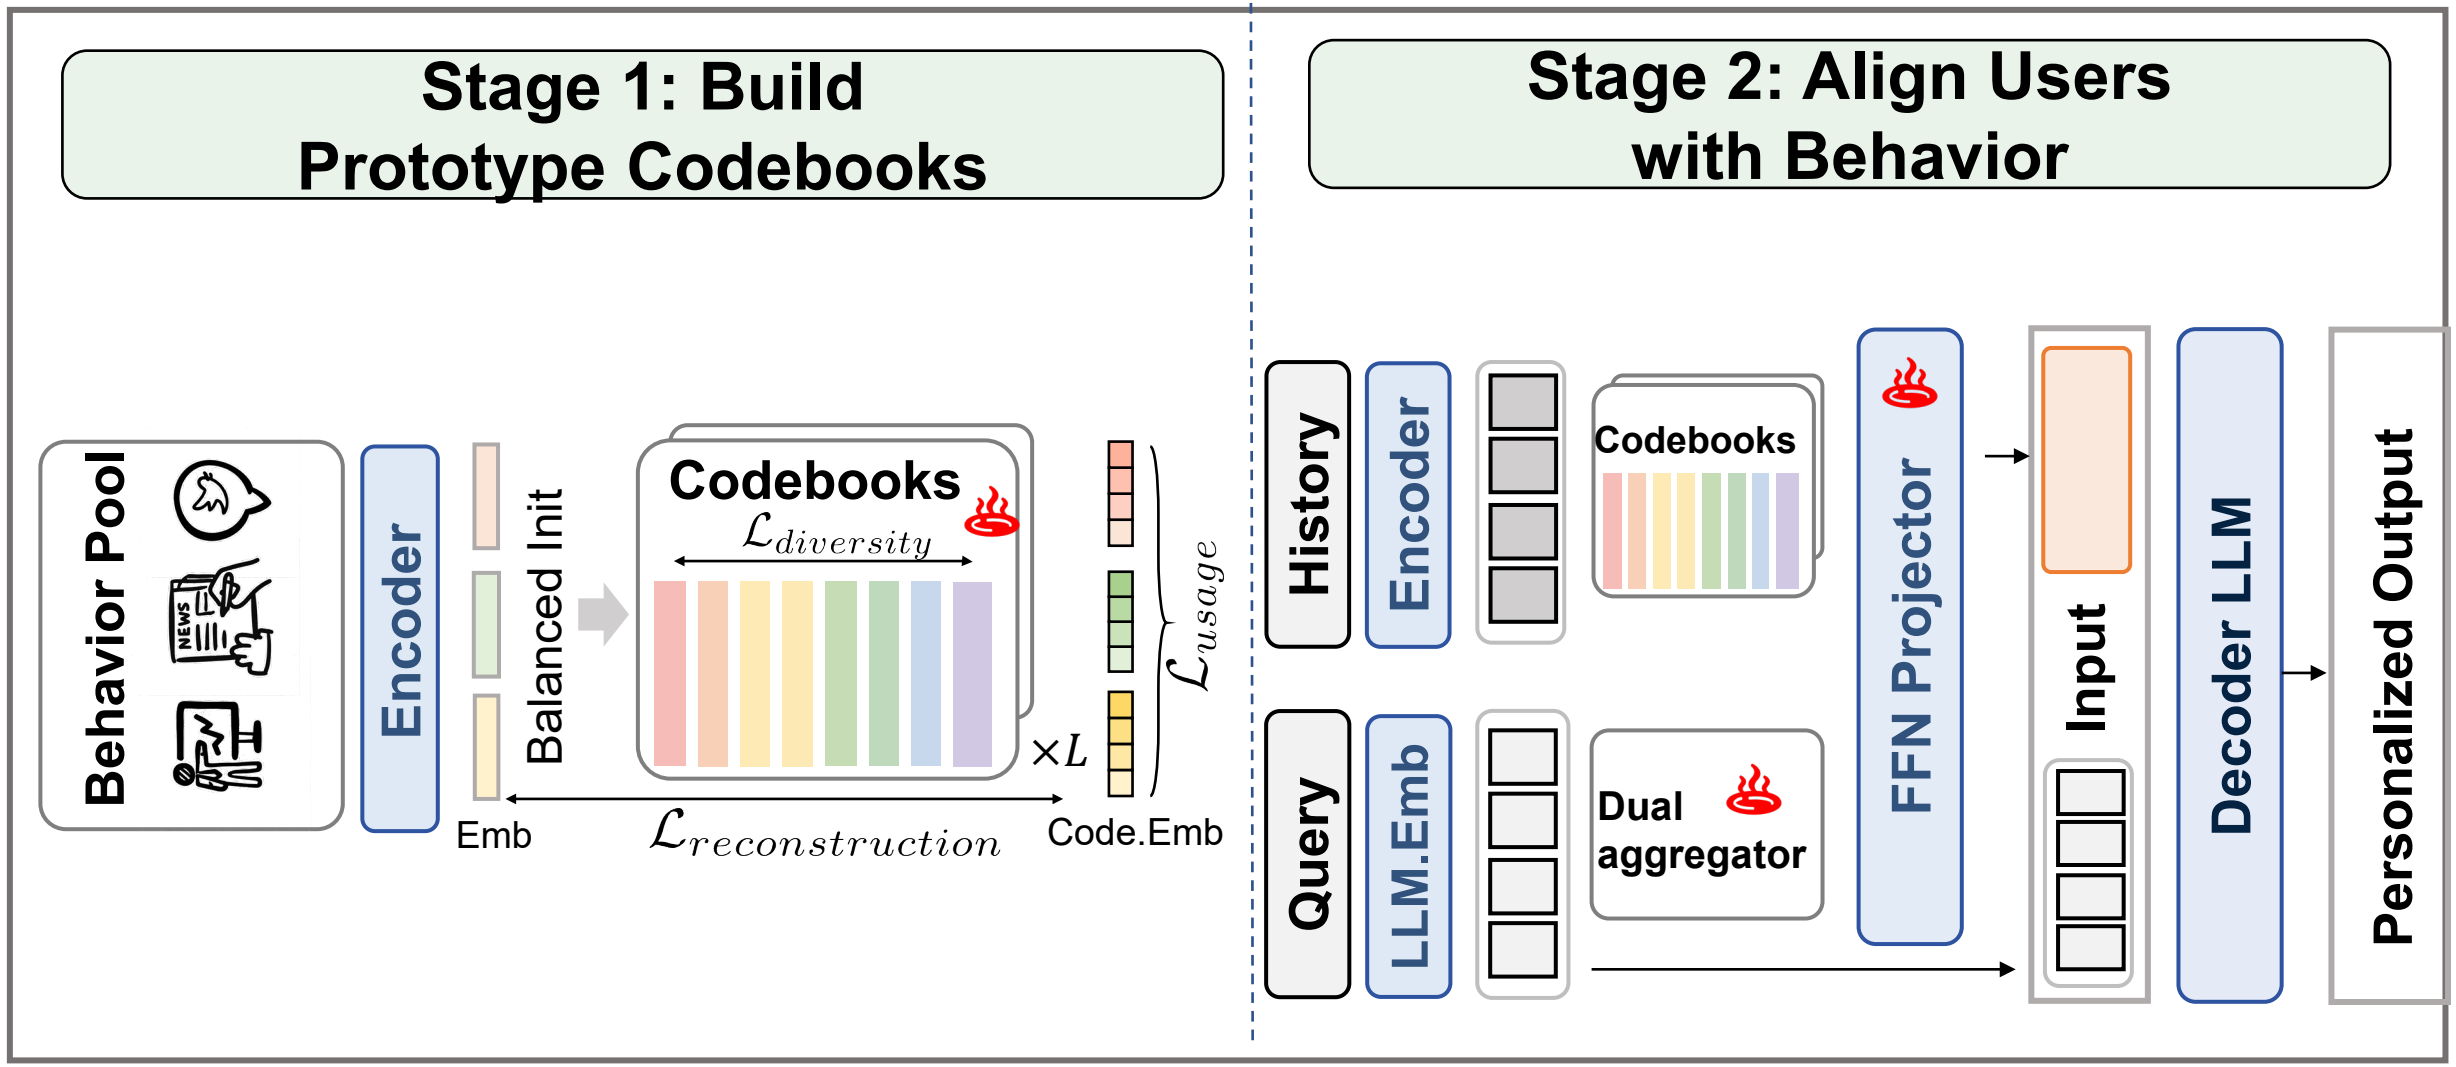
- **Stable patterns** capture enduring preferences.
- The **current query** determines which history matters
- Generation combines **stable identity** and **dynamic context.**



Stable identity + **Dynamic context** → **Personalized response**

Personalization stays consistent over time and adapts to context.

Training Framework: Two Stages



CURP: compact models, strong results

≈28M

trainable parameters

>70%

of LaMP metrics reach the best reported score

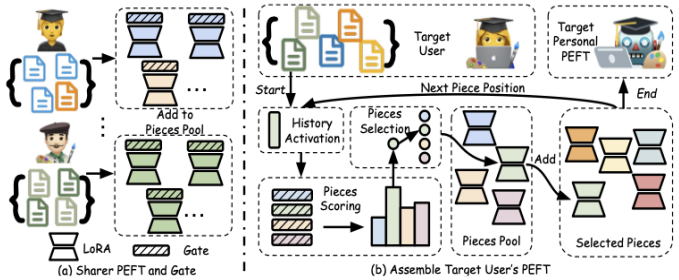
Both representations contribute

LaMP-2 variant	Accuracy
Full model	0.629
Without stable representation	0.539
Without dynamic representation	0.474

Wang et al., CURP, EMNLP 2026.

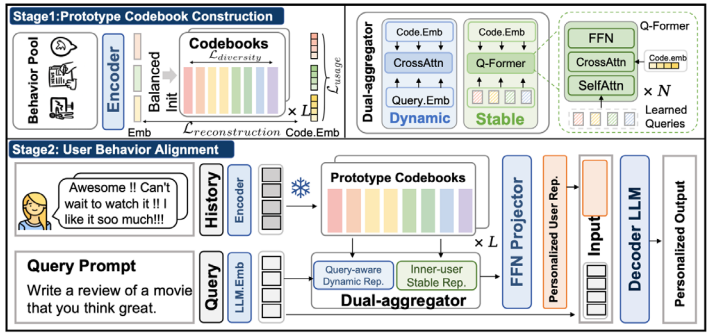
Personalized LLMs: Internal and External Control

Router / MoE
Internal parameter activation



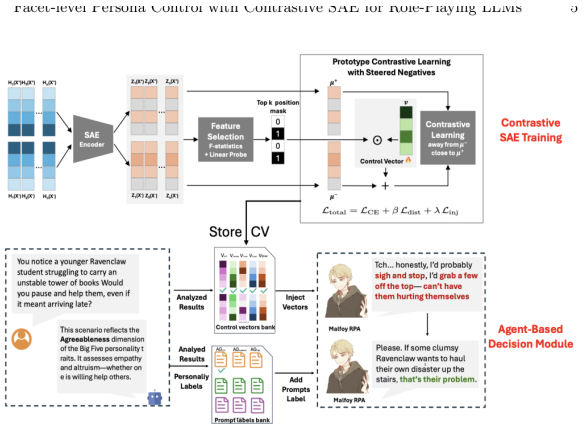
Share parameters
to lower per-agent cost

Soft Prompt
External guidance



Lightweight guidance
with a frozen base model

SAE / Activation Steering
Representation guidance

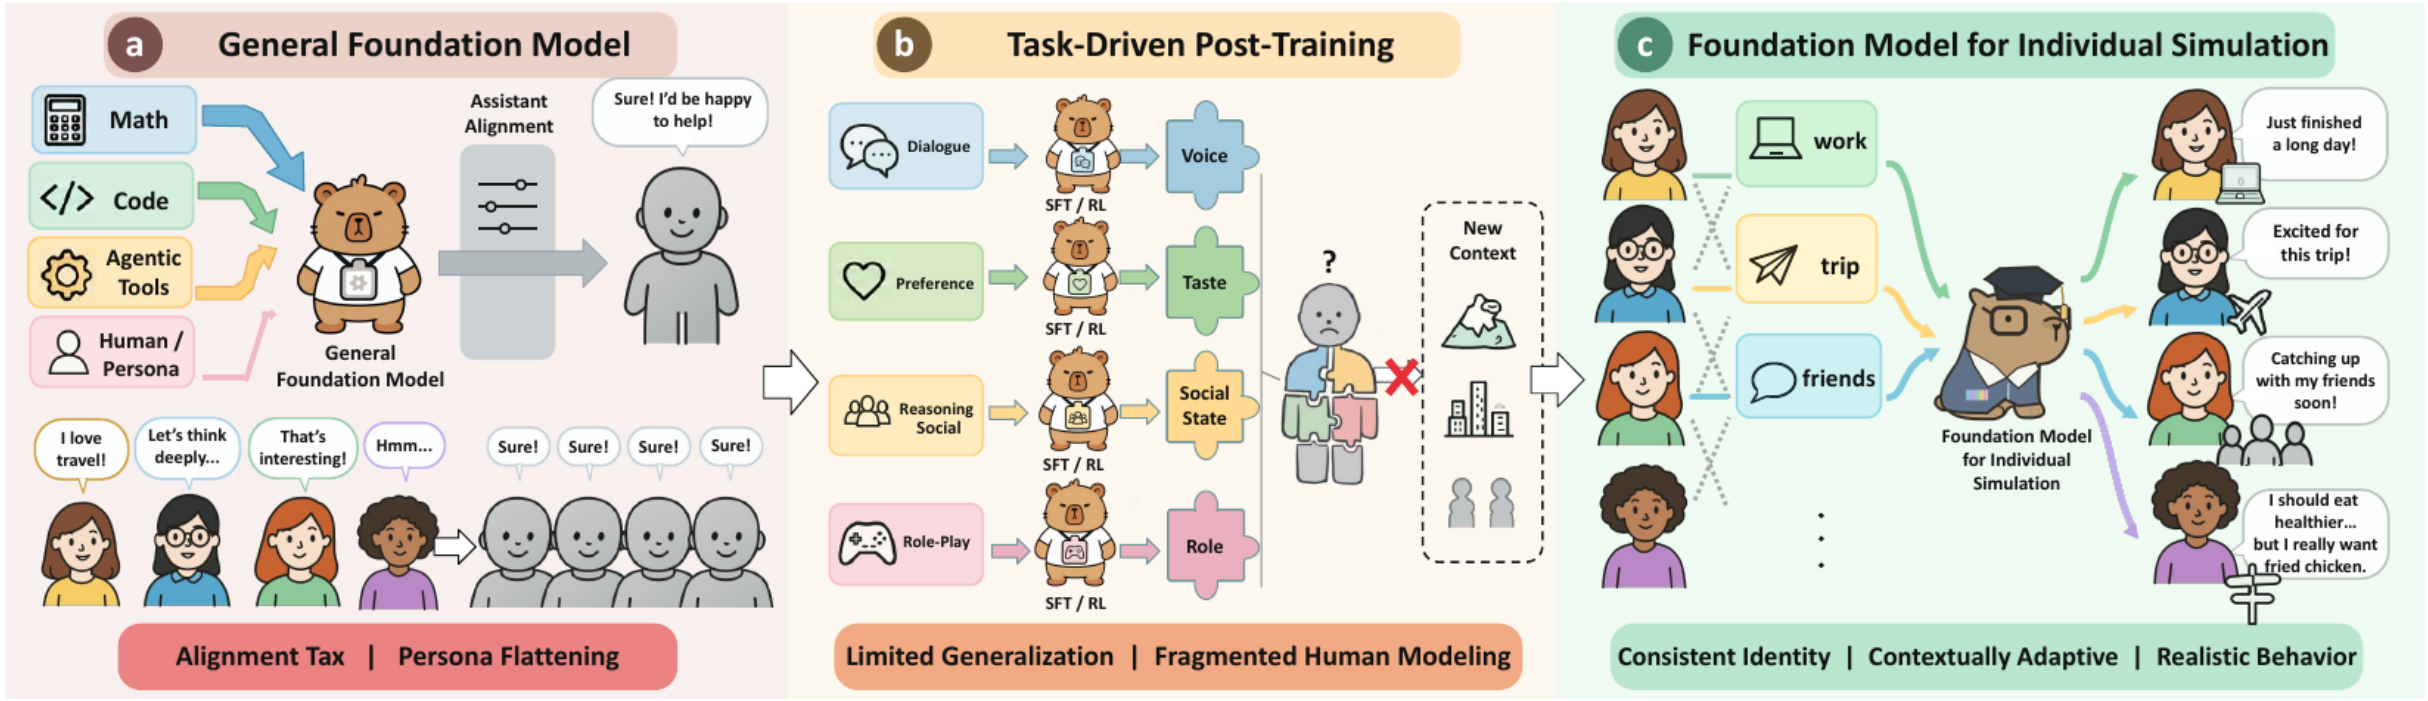


Explain or control
persona and bias directions

The next generation of personalized parameters should combine internal control with external guidance.

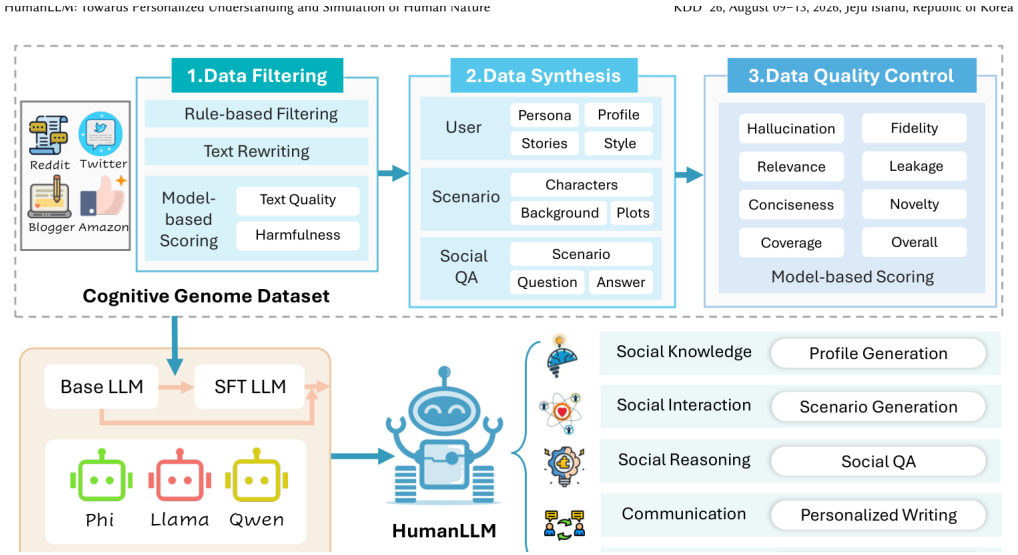
Mixture-of-Personas Language Models, 2025; Tang et al., Contrastive SAE Persona Control, 2026.

Why task skill is not yet a person model

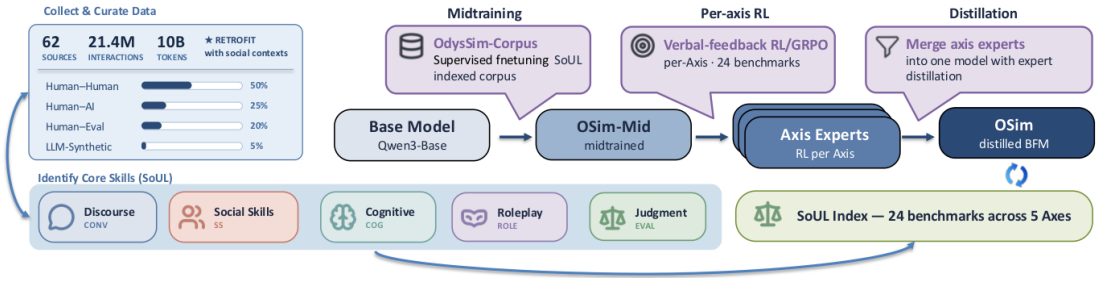


What the capability should the model learn, so that personalization works across tasks and contexts?

Recent Attempts at Human Behavior Foundation Models



HumanLLM KDD26



OdysSim Neurips26

The **prediction of individual preferences, thoughts, and writing styles**.

A **broad range of social-behavior tasks** and performs well across settings.

The capability landscape is not well defined.

Lei et al., HumanLLM, KDD 2026; Zhou et al., OdysSim, NeurIPS 2026.

Socio-Foundation

FONTS

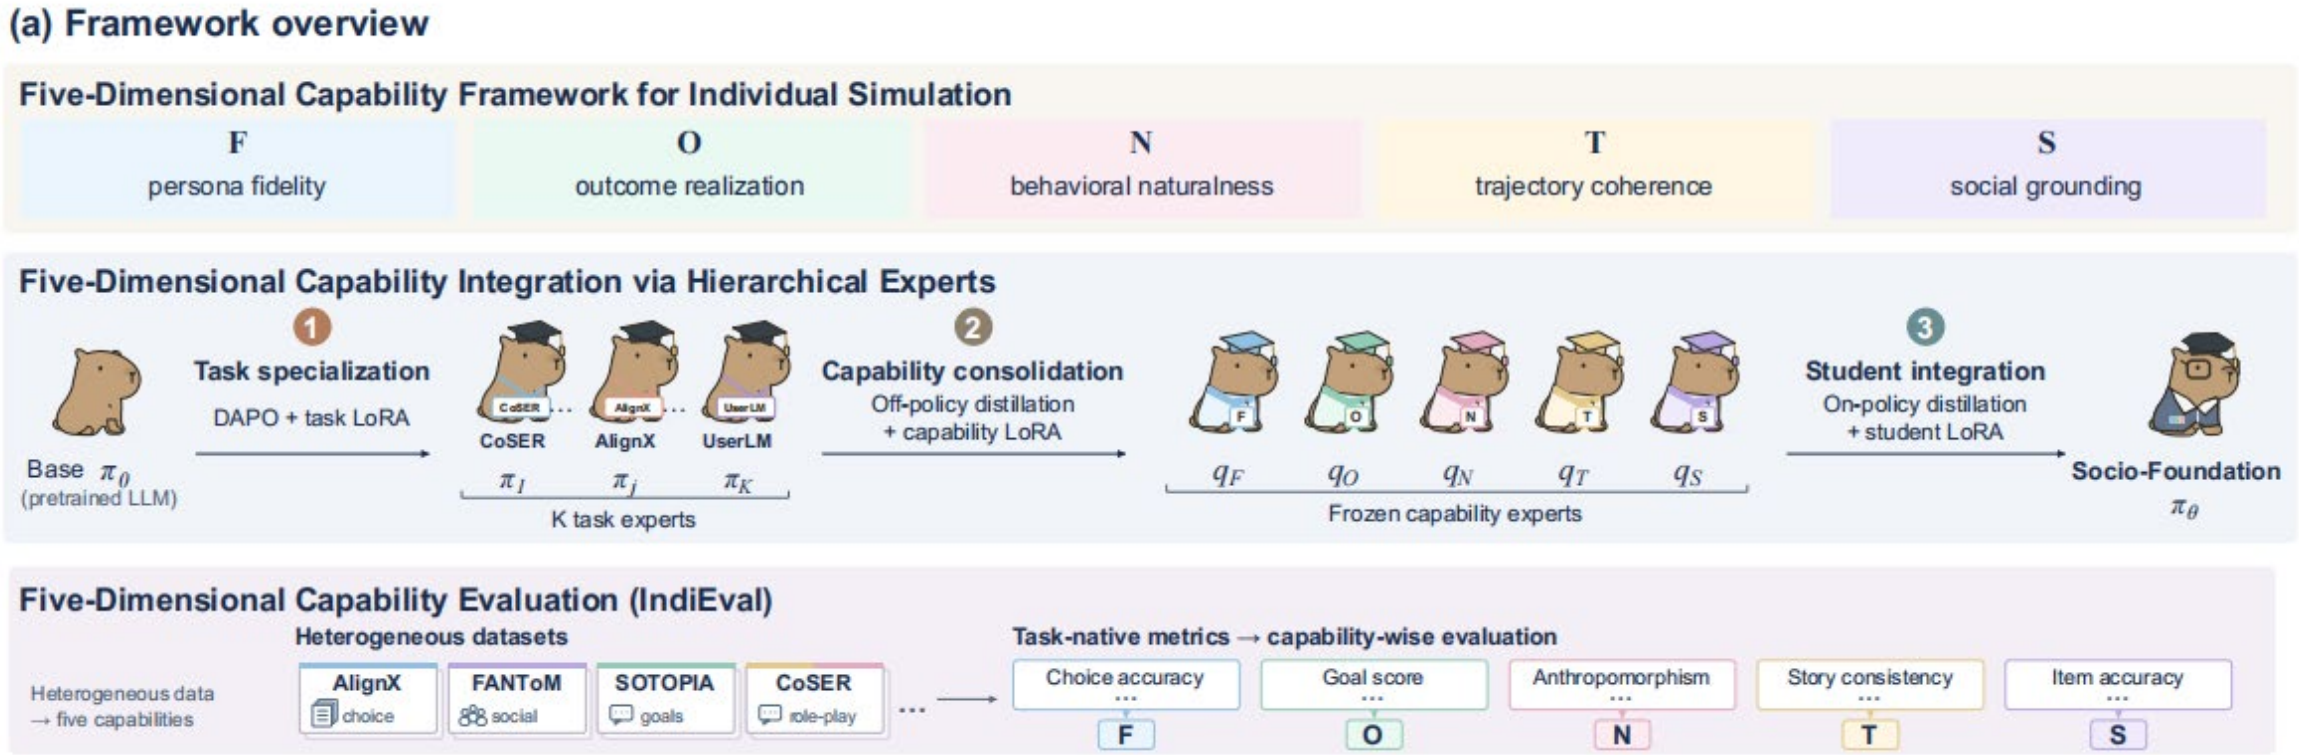
Five behavioral capabilities

Training

Task and capability experts

IndiEval

29 evaluation metrics



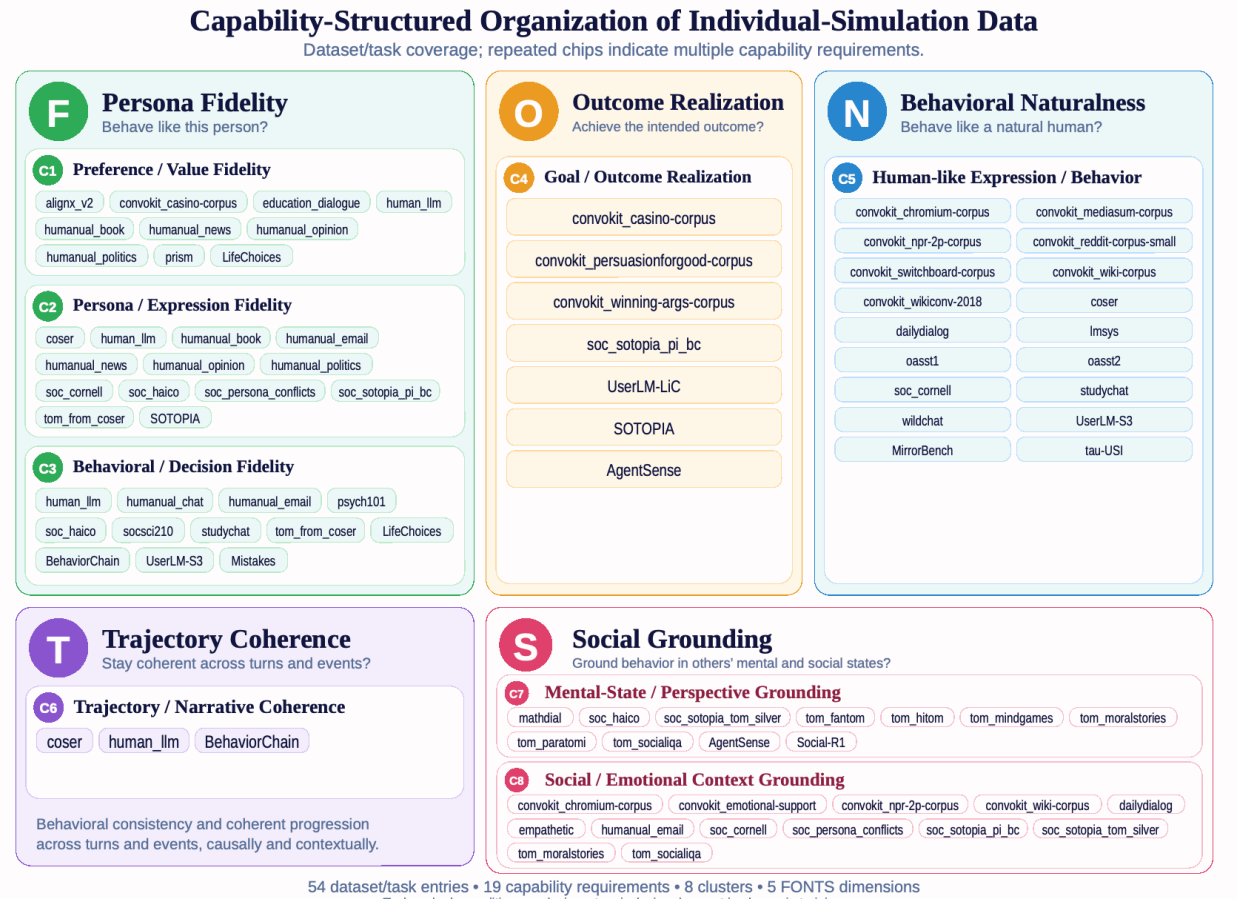
Socio-Foundation: A Model for Generalizable Individual Behavior Simulation via Hierarchical Capability Distillation, Arxiv'26

FONTS: five dimensions of individual simulation

F	**Persona fidelity**	Preferences, values, and decisions
O	**Outcome realization**	Goals and intended outcomes
N	**Behavioral naturalness**	Human-like expression and interaction
T	**Trajectory coherence**	Continuity over time
S	**Social grounding**	Beliefs and private information

How the FONTS taxonomy was derived

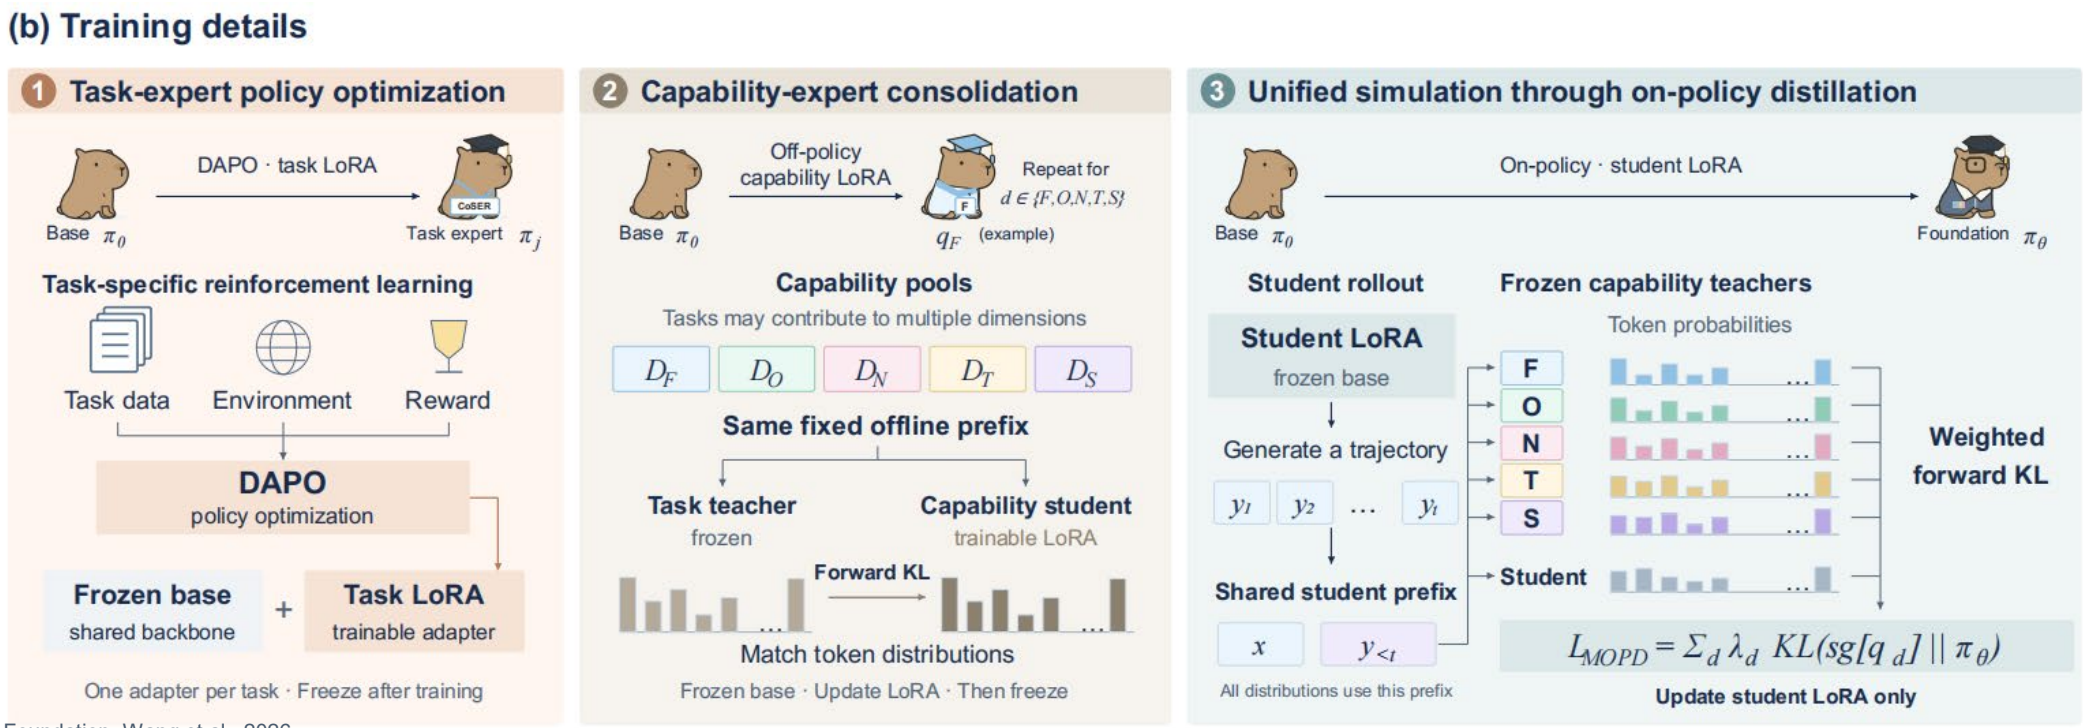
- Analyze behavior requirements for each **dataset–task–role** setting.
- Annotate **19 requirements** and group them into **8 clusters.**
- McAdams & Pals' New Big Five guides **five FONTS dimensions**.
- Tasks can span dimensions when they require different behaviors.



Bottom-up task analysis, followed by theory-guided consolidation.

Socio-Foundation: Three-stage Training

- Task experts: one task LoRA per task, trained **with DAPO.**
- Capability experts: group task experts via **off-policy distillation (MOPD).**
- Unified simulator: integrate five capability experts **via MOPD.**



IndiEval: Five Behavioral Capabilities

5 capabilities / 11 targets / 29 metrics / 6,297 cases

Capability	Example evaluation target	Example metric
Persona fidelity	Preference and decision fidelity	Choice accuracy
Outcome realization	Goal accomplishment	Goal score
Behavioral naturalness	Human-like expression	Anthropomorphism
Trajectory coherence	Narrative continuity	Story consistency
Social grounding	Mental-state inference	Item accuracy

13 benchmarks and 4 additional out-of-distribution tests

Socio-Foundation, Wang et al., 2026.

Socio-Foundation: Performance on IndiEval

+11.0 points over the Qwen3-8B base model

	F									O							N								T			S			Avg.
Model	LifeChoices Acc. ↑	AlignX	Choice acc. ↑	HumanLLM Hit@5 ↑	BehaviorChain Node ↑	CoSER Char. fid. ↑	HUMANUAL Response ↑	HUMANUAL State ↑	UserLM-S3 Intent adh. ↑	SOTOPIA Believ. ↑	SOTOPIA Goal ↑	SOTOPIA Relation. ↑	SOTOPIA Knowledge ↑	SOTOPIA Benefits ↑	AgentSense Goal maj. ↑	UserLM-LiC Intent cov. ↑	UserLM-LiC Task score ↑	UserLM-S3 Term. F1 ↑	MirrorBench GTEval ↑	CoSER Anthrop. ↑	UserLM-S3 Human lik. ↑	UserLM-S3 Diversity ↑	UserLM-S3 Inv. overlap ↑	UserLM-S3 Role adh. ↑	CoSER Story cons. ↑	BehaviorChain CumScore ↑	CoSER Story qual. ↑	FANToM Item acc. ↑	Social-R1 Acc. ↑	AgentSense Private acc. ↑	Average ↑
Frontier models																															
GPT-5.5	**92.0**	61.0	**64.1**	**89.7**	**49.6**	**21.8**	44.3	98.0	84.7	**59.0**	63.7	29.3	**55.0**	**96.7**	**99.0**	**94.4**	0.0	41.3	41.9	51.2	91.6	75.6	32.0	62.6	**58.7**	65.6	79.3	**69.0**	82.4	**67.2**
Open-weight general-purpose models																															
DeepSeek-V4-Flash	81.8	59.6	62.0	69.3	43.2	15.5	39.5	96.0	82.7	54.5	60.8	27.4	51.9	95.5	93.3	57.9	34.1	48.5	**44.5**	29.8	86.8	87.5	5.0	62.6	26.3	**75.7**	61.5	51.0	82.4	59.7
GLM-5.2	84.2	62.2	63.4	76.7	43.9	20.0	**45.5**	94.0	86.1	55.5	**72.9**	**37.2**	50.5	**96.7**	97.3	76.8	0.0	43.2	40.8	**79.7**	91.0	78.4	6.0	59.7	34.1	68.0	74.3	**69.0**	80.8	63.2
Kimi-K2.6	77.4	**63.8**	51.8	82.9	27.3	18.2	42.7	98.0	**86.8**	51.6	67.9	35.5	49.6	95.9	95.5	65.2	30.3	37.6	30.4	58.9	85.8	82.9	21.0	52.3	42.7	53.3	75.5	63.0	**91.4**	61.2
Qwen3-Next-80B	76.0	58.4	48.6	45.8	13.3	16.6	41.8	86.0	83.3	51.0	65.1	29.2	48.9	92.4	92.2	50.7	0.0	38.5	18.8	27.5	87.9	85.9	23.0	39.8	11.2	38.2	71.1	57.0	80.8	50.9
Qwen3-235B-A22B	76.0	59.2	56.8	61.3	22.6	17.3	42.3	83.0	84.4	52.8	67.7	26.2	51.5	95.3	92.2	54.8	27.6	45.6	25.1	48.7	93.0	78.0	20.0	49.3	18.6	51.9	71.6	60.0	79.4	55.6
Specialized simulation models																															
OSim 8B	80.4	60.0	60.2	88.8	9.0	12.6	40.1	98.0	48.9	32.2	45.8	26.5	46.2	90.2	68.9	34.8	0.0	51.7	8.6	59.3	**94.4**	37.8	93.0	44.7	57.8	28.7	65.5	62.0	87.7	55.4
CoSER 8B	60.2	55.0	12.2	33.0	31.0	3.4	11.1	97.0	45.9	27.5	44.1	17.2	47.8	60.6	17.98	8.9	42.6	11.7	31.1	12.2	85.2	36.9	82.0	55.5	5.9	47.6	52.6	47.0	57.0	38.2
HumanLM 8B	67.8	55.6	58.0	33.9	18.3	15.6	39.3	83.0	77.3	44.9	56.1	25.2	49.7	87.6	74.8	35.7	6.5	38.6	23.4	24.4	89.2	73.0	72.0	53.4	6.4	41.5	62.4	48.0	78.2	48.6
UserLM 8B	1.6	12.6	0.2	0.1	44.2	2.3	12.6	**100.0**	28.7	19.8	31.4	5.0	50.1	1.4	9.1	2.9	**57.8**	4.8	25.5	54.4	49.9	56.2	**100.0**	**77.4**	0.0	42.6	10.5	5.0	16.6	21.5
USP 8B	0.0	0.6	0.0	0.0	18.4	1.6	10.6	94.0	21.0	19.5	28.8	2.3	47.7	10.0	0.0	0.0	0.0	23.6	16.9	65.3	50.3	76.6	99.0	63.9	0.0	37.8	2.4	1.0	24.7	19.3
Base model and Socio-Foundation																															
Qwen3-8B	72.2	57.0	56.4	34.9	21.0	14.6	38.7	93.0	77.5	46.4	54.2	27.1	51.6	90.8	70.7	29.3	12.9	40.1	25.5	25.4	87.4	75.9	70.0	54.0	6.9	44.3	75.2	50.0	80.5	50.7
Socio-Foundation	80.0	61.0	58.2	86.40	25.0	13.4	39.5	**100.0**	74.4	45.2	54.2	26.5	52.9	93.1	72.8	34.7	52.5	**54.1**	31.6	47.9	90.2	**97.1**	**100.0**	54.7	49.35	48.6	**79.6**	64.0	83.7	61.7

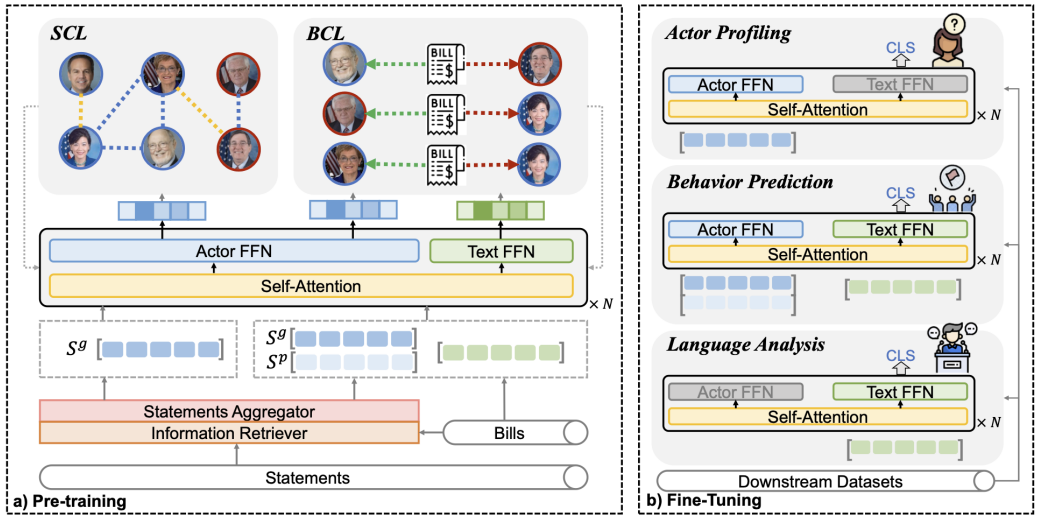
Beyond Simulation: Individual Decision Intelligence

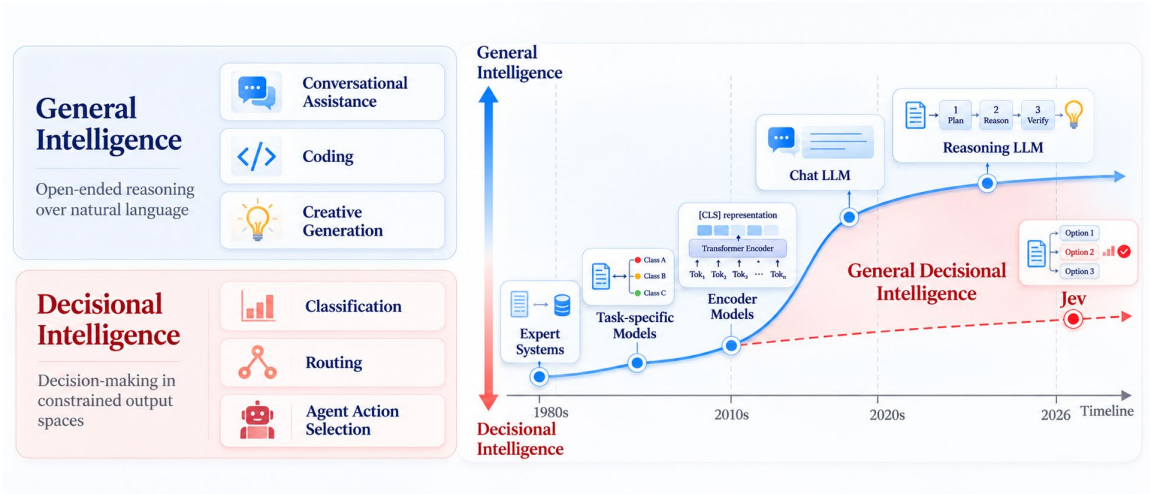
Learned decisions

Model the person's choices

Agent-based simulation

Model interactions and collective dynamics





Decision models could provide learned behavioral rules for agents

Mou et al., UPPAM, ACL 2023; Lin et al., Disc-Jev, 2026; TypeSafe AI, Jev, 2026

Social simulation: individual survey prediction via Inner-Jev

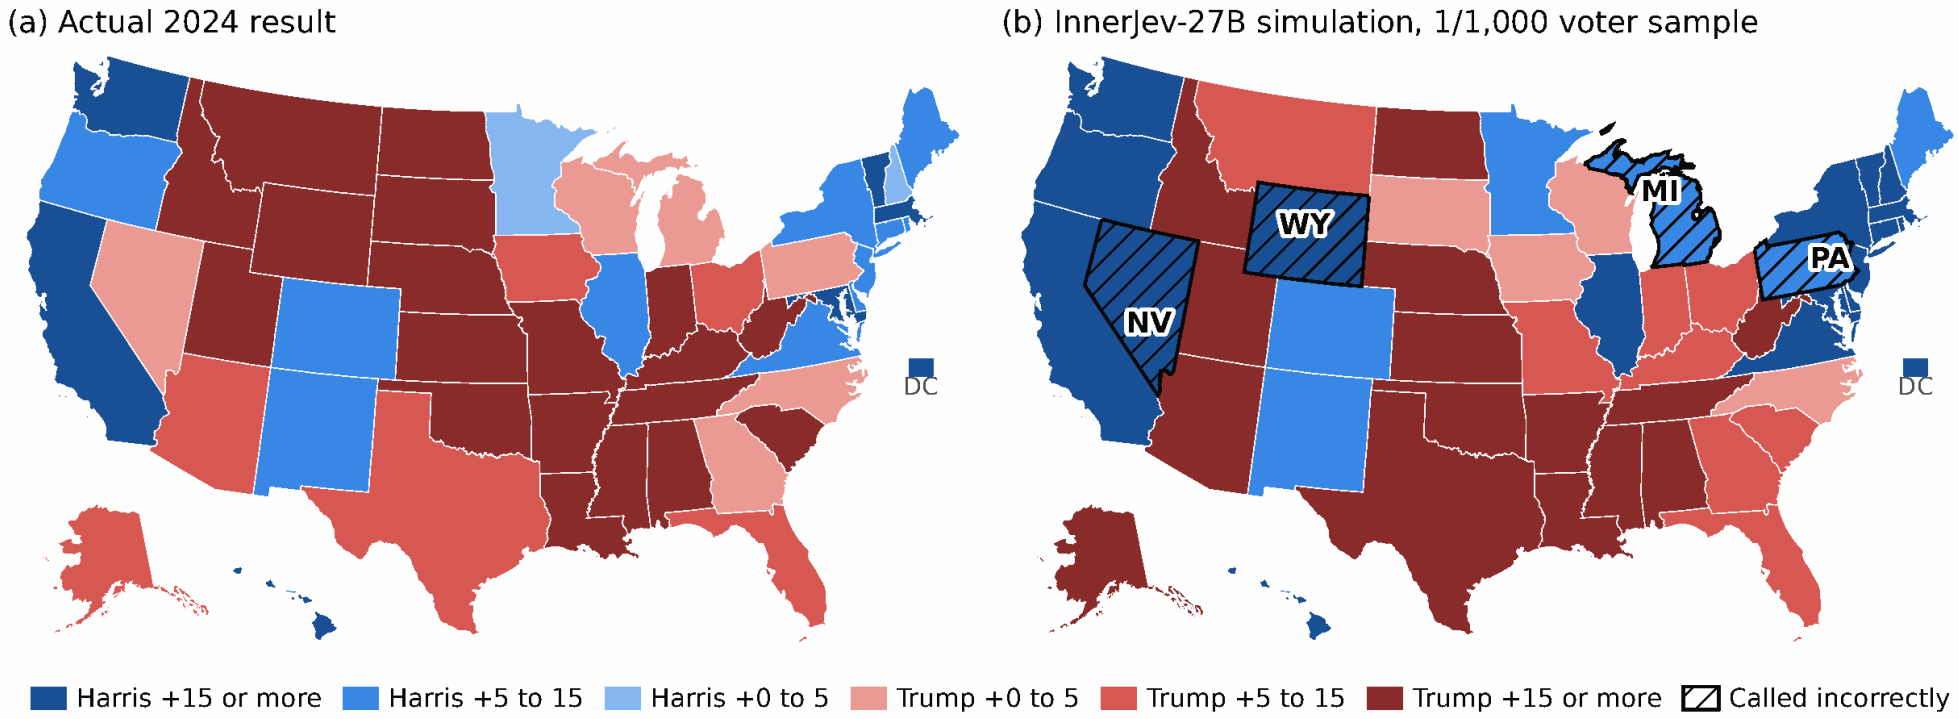
Finding: Decision models predict individual survey responses on par with generative LLMs at a much lower cost.

277,500 respondent-question pairs across 10 domains. ACC (%), latency (s).

	Jev		Qwen3.5-4B				DeepSeek-V4-Flash-0731				Ours: InnerJev-27B	
	Native		Non-thinking		Thinking		Non-thinking		Thinking (high)		Native	
Domain	**ACC ↑**	**Latency ↓**	**ACC ↑**	**Latency ↓**	**ACC ↑**	**Latency ↓**	**ACC ↑**	**Latency ↓**	**ACC ↑**	**Latency ↓**	**ACC ↑**	**Latency ↓**
Citizenship	43.00	0.478	39.46	6.999	41.84	102.333	43.60	268.818	**45.41**	302.180	45.09	8.412
Environment	34.85	0.480	30.30	6.832	33.62	101.452	35.51	234.424	**38.29**	283.856	36.21	8.919
Family	37.23	0.493	29.17	6.728	33.55	106.760	40.31	260.248	**42.94**	328.831	40.77	8.490
Health	35.44	0.494	30.46	6.816	35.03	97.687	34.83	250.956	37.44	292.322	**39.99**	8.623
National Identity	37.70	0.504	31.69	7.473	34.30	98.672	35.93	249.620	37.92	182.197	**38.86**	8.618
Religion	39.69	0.494	34.77	8.319	38.39	101.512	39.90	81.061	**42.42**	178.482	40.65	6.131
Role of Government	38.66	0.519	34.97	7.513	36.17	97.840	38.02	74.523	**39.86**	189.612	38.39	8.321
Social Inequality	36.03	0.486	29.05	7.115	31.34	103.303	34.83	104.599	36.75	147.013	**36.82**	8.711
Social Networks	36.82	0.497	34.53	7.117	37.72	110.469	37.54	112.630	**39.96**	115.594	38.51	8.898
Work Orientations	36.27	0.506	32.36	6.164	36.70	113.625	37.33	113.121	41.93	111.529	**42.39**	8.668
Overall	37.69	0.496	32.87	7.124	36.04	103.394	37.90	173.999	**40.41**	211.649	39.87	8.357

Social simulation: population election outcomes via Inner-Jev

Finding: Affordable population simulation still inherits partisan bias.

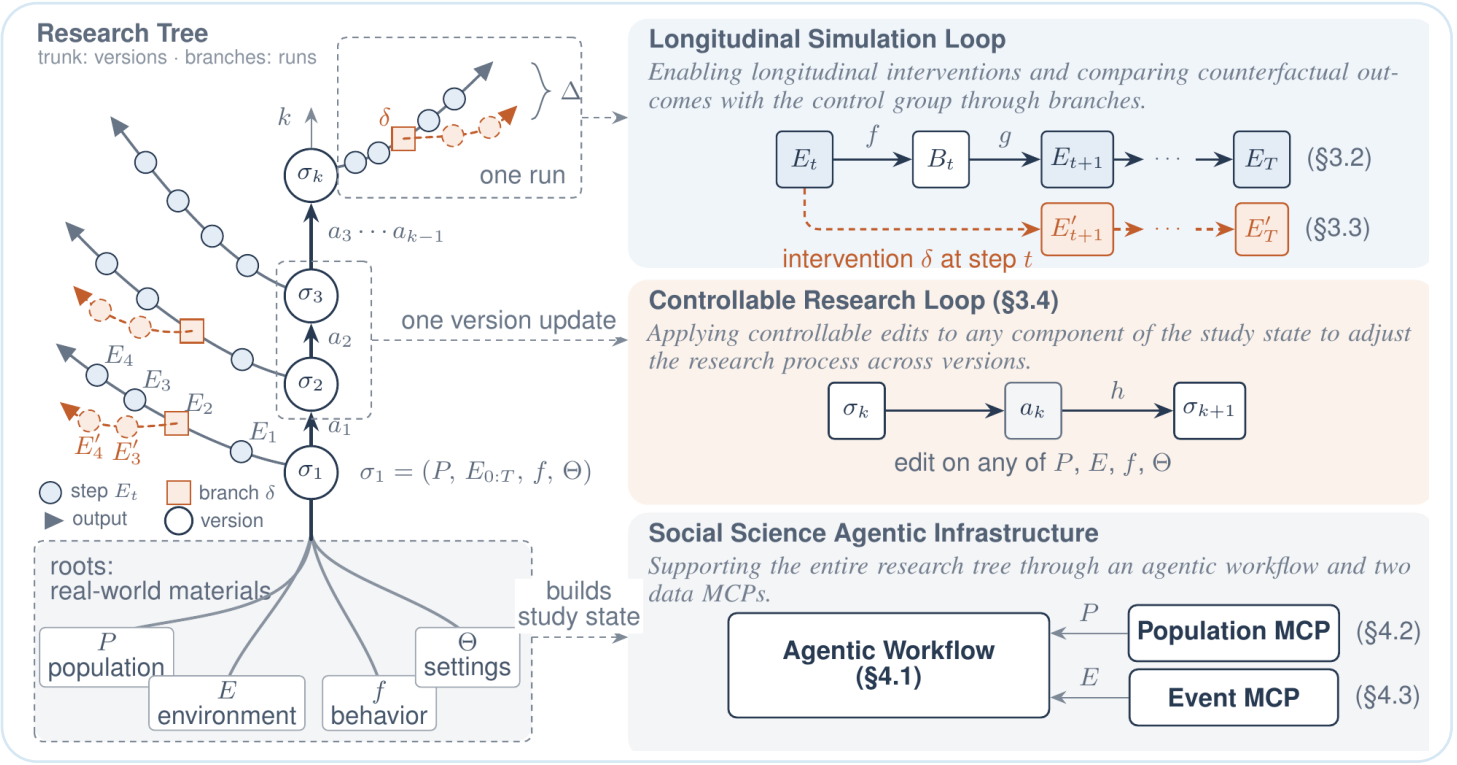


Voter model	States ↑	Battleground ↑	MAE ↓	RMSE ↓	Bias
ElectionSim (GPT-4o-mini)	46/51	11/15	**3.06**	**3.80**	**+2.31**
InnerJev-27B (decision model)	**47/51**	**12/15**	6.52	9.42	+3.73

Duan et al., General Decision Models, 2026.

SocioVerse2: Two Loops and One Infrastructure

Longitudinal social simulation under a human–AI co-evolutionary paradigm



BRANCHES · RUNS & INTERVENTIONS δ

Longitudinal Simulation Loop

Evolve · intervene · branch: counterfactual control groups that reality cannot supply

TRUNK · VERSIONS OF THE STUDY

Controllable Research Loop

The study is an editable state; researchers author edits to P, E, f, Θ at every stage

ROOTS · REAL-WORLD P, E, f, Θ

Social Science Agentic Infrastructure

7 skills with human checkpoints; Population MCP (~1B personas) + Event MCP (21 sources)

From cross-sectional prediction to longitudinal trajectory simulation.

Zhang et al., SocioVerse2: A Longitudinal Dynamic Social Simulation Framework under a Human–AI Co-evolutionary Paradigm, arXiv:2609.24911, 2026.

SocioVerse Challenge 2026

AI4SS Challenge for Human–AI Collaboration and Social Governance

A TRACK A

Historical Reconstruction & Forecasting

Rebuild the past, predict the next

B TRACK B

Counterfactual Intervention & Impact Assessment

What if a policy changes?

C TRACK C

Methodological & Systems Innovation

Better models, agents and tools

$17,600

total prizes & support

18 main awards · $12,600

1st $1,500 × 3
2nd $750 × 6
3rd $400 × 9

$5,000 API & compute credits

KEY DATES

Sep 15	Sep 25	Oct 12	Oct 15	Oct 31	Nov 1–7	Nov 8
Launch	Code v0.2.0	**Proposal**	Screening	**Final submission**	Review	Awards, Rome

Teams of 1–5

Online or local

Your own topic

Scan to register

socioverse.fudan-disc.com/challenge

Proposals due Oct 13 · Final submission Oct 31 · Awards at CIKM 2026, Rome

SocioVerse Challenge 2026: Organizers & Committees

PLATFORM SUPPORT

上海创智学院
Shanghai Innovation Institute

ORGANIZERS

DISC Lab, Fudan University

NLP Group, King's College London

Behavioral Science Center, Fudan University

Center for Computational and Intelligent Communication, Fudan University

EVALUATION COMMITTEE



Xinnong Zhang
SII / Fudan



Jiayu Lin
SII / Fudan



Jia Wang
SII / Tongji



Tao Chiang
Fudan



Hanjia Lyu
SMU



Hanqi Yan
King's College London



Siyuan Wang
CUHK



Yunlu Yin
Fudan

ADVISORY COMMITTEE



Zhongyu Wei
SII / Fudan



Baohua Zhou
Fudan



Liyin Jin
Fudan



Xuanjing Huang
SII / Fudan



Yulan He
King's College London



Jiebo Luo
Rochester



Liang Pang
ICT, CAS



Guanying Li
Fudan



Rongwei Chu
Fudan



Lingyi Zhou
Fudan



Siming Chen
Fudan



Zhenfei Yin
Oxford

Awards announced at LASS 2026 @ CIKM 2026, Rome · socioverse.fudan-disc.com/challenge

Source: socioverse.fudan-disc.com/challenge/#committee

Personalized Foundation Model for Social Simulation

Wei, Zhongyu

Fudan University | Shanghai Innovation Institute

2026.10.09

Second Workshop on Social Simulation with LLMS: Fidelity in Applications

@ COLM 2026, Hilton Union Square, San Francisco, CA, USA